

VISVESVARAYA TECHNOLOGICAL UNIVERSITY

“JnanaSangama”, Belgaum -590014, Karnataka.



LAB REPORT on

COURSE TITLE

Submitted by

S Swaroop Bharadwaj

in partial fulfillment for the award of the degree of
BACHELOR OF ENGINEERING
in
COMPUTER SCIENCE AND ENGINEERING



B.M.S. COLLEGE OF ENGINEERING

(Autonomous Institution under VTU)

BENGALURU-560019

May-2023 to July-2023

B. M. S. College of Engineering,
Bull Temple Road, Bangalore 560019
(Affiliated To Visvesvaraya Technological University, Belgaum)
Department of Computer Science and Engineering



CERTIFICATE

This is to certify that the Lab work entitled “**BIG DATA ANALYTICS**” carried out by **S Swaroop Bharadwaj (1BM20CS136)**, who is bonafide student of **B. M. S. College of Engineering**. It is in partial fulfillment for the award of **Bachelor of Engineering in Computer Science and Engineering** of the Visvesvaraya Technological University, Belgaum during the year 2022. The Lab report has been approved as it satisfies the academic requirements in respect of a **Big Data Analytics - (20CS6PEBDA)** work prescribed for the said degree.

Vikranth BM
Assistant Professor
Department of CSE
BMSCE, Bengaluru

Dr. Jyothi S Nayak
Professor and Head
Department of CSE
BMSCE, Bengaluru

Index Sheet

Sl.No.	Experiment Title
1	Cassandra Lab Program1: - Employee Database
2	Cassandra Lab Program1: - Library Database
3	MongoDB- CRUD Demonstration
4	Hadoop installation
5	Hadoop Commands
6	Hadoop Program: Average Temperature
7	Hadoop Program: Word Count
8	Hadoop program: Join operation
9	Scal Program
10	Scala Program: Word Count

Course Outcome

CO1	Apply the concept of NoSQL, Hadoop or Spark for a given task
CO2	Analyze the Big Data and obtain insight using data analytics mechanisms.
CO3	Design and implement big data applications by applying NoSQL, Hadoop or Spark

1 Perform the following DB operations using Cassandra.

1. Create a keyspace by name Employee

2. Create a column family by

name Employee-Info with

attributes Emp_Id Primary Key,

Emp_Name,

Designation, Date_of_Joining, Salary, Dept_Name

3. Insert the values into the table in batch

4. Update Employee name and Department of Emp-Id 121

5. Sort the details of Employee records based on salary

6. Alter the schema of the table Employee_Info to add a column Projects which stores a set of

Projects done by the corresponding Employee.

7. Update the altered table to add project names.

8. Create a TTL of 15 seconds to display the values of Employees.

```
cqlsh:employee> CREATE KEYSPACE employee WITH REPLICATION={ 'class' : 'SimpleStrategy', 'replication_factor' : 1};
```

```
cqlsh:employee> USE employee;
```

```
cqlsh:employee> create table employee_info(emp_id int PRIMARY KEY, emp_name text, ... designation text, date_of_joining timestamp, salary double PRIMARY KEY, dept_name text);
```

```
cqlsh:employee> CREATE TABLE employee_info(emp_id int, emp_name text, designation text, date_of_joining timestamp, salary double, dept_name text, PRIMARY KEY(emp_id, salary));
```

```
cqlsh:employee> BEGIN BATCH INSERT INTO
```

...

```

employee_info(emp_id,emp_name,designation,date_of_joining,sal
ary,dept_name) ...
VALUES(100,'John','MANAGER','2021-09-11',30000,'TESTING');
... INSERT INTO
...
employee_info(emp_id,emp_name,designation,date_of_joining,sal
ary,dept_name) ...
VALUES(111,'Tom','ASSOCIATE','2021-06-22',25000,'DEVELOPING');
... INSERT INTO
...
employee_info(emp_id,emp_name,designation,date_of_joining,sal
ary,dept_name) ...
VALUES(121,'Elsa','MANAGER','2021-03-30',35000,'HR');
... INSERT INTO
...
employee_info(emp_id,emp_name,designation,date_of_joining,sal
ary,dept_name) ...
VALUES(115,'Chris','ASSISTANT','2021-12-30',20000,'DEVELOPING');
... INSERT INTO
...
employee_info(emp_id,emp_name,designation,date_of_joining,sal
ary,dept_name) ...
VALUES(105,'Sarah','ASSOCIATE','2021-06-25',25000,'TESTING');
... APPLY BATCH;
cqlsh:employee> SELECT * FROM employee_info
... ;

```

```
cqlsh:employee> UPDATE employee_info SET emp_name = 'Jessica', dept_name =  
'DEVELOPING' WHERE emp_id = 121;
```

```
cqlsh:employee> UPDATE employee_info SET emp_name = 'Jessica', dept_name =  
'DEVELOPING' WHERE emp_id = 121 AND salary = 35000;
```

```
cqlsh:employee> SELECT * FROM employee_info ;  
MANAGER | John
```

```
cqlsh:employee> SELECT * FROM employee_info WHERE emp_id in (105, 111, 121, 115,  
100) order by salary; cqlsh:employee> paging off
```

Disabled Query paging.

```
cqlsh:employee> SELECT * FROM employee_info WHERE emp_id in (105, 111, 121, 115, 100)  
order by salary;
```

```
cqlsh:employee> ALTER TABLE employee_info ADD projects text;
```

```
cqlsh:employee> UPDATE employee_info SET projects = 'Chat App' WHERE emp_id = 111;
```

```
cqlsh:employee> UPDATE employee_info SET projects = 'Chat App' WHERE emp_id = 111  
and salary = 25000;
```

```
cqlsh:employee> UPDATE employee_info SET projects = 'Discord Bot' WHERE emp_id  
= 115 and salary = 20000;
```

```
cqlsh:employee> UPDATE employee_info SET projects = 'Campus Portal' WHERE emp_id  
= 105 and salary = 25000;
```

```
cqlsh:employee> UPDATE employee_info SET projects = 'YouTube Downloader' WHERE  
emp_id = 100 and salary = 30000;
```

```
cqlsh:employee> UPDATE employee_info SET projects = 'Library Management System '  
WHERE emp_id = 121 and salary = 35000;
```

```
cqlsh:employee> SELECT * FROM employee_infor  
... ;
```

```
cqlsh:employee> SELECT * FROM employee_info ;
```

```
cqlsh:employee> INSERT INTO
```

...

employee_info(emp_id,emp_name,designation,date_of_joining,salary,dept_name) ...

... ;

cqlsh:employee> INSERT INTO

...

employee_info(emp_id,emp_name,designation,date_of_joining,salary,dept_name) ...

VALUES(110,'SAM','ASSOCIATE','2021-01-11',28000,'TESTING')

USING TTL 15;

cqlsh:employee> SELECT TTL(emp_name) from employee_info

WHERE emp_id = 110; ttl(emp_name)

3

cqlsh:employee> SELECT * FROM employee_info;

Output:

```
Activities Terminal Apr 29 10:43 bmsce@bmsce-Precision-T1700 ~
cqlsh:employee> insert into EmployeeInfo (emp_name,emp_ID,designation,date_of_joining,salary,dept_name) values('jal',121,'manager','2022-03-13',40000,'cs');
cqlsh:employee> APPLY BATCH
... Insert into EmployeeInfo (emp_name,emp_ID,designation,date_of_joining,salary,dept_name) values('jal',121,'manager','2022-03-13',40000,'cs');
cqlsh:employee> BEGIN BATCH insert into EmployeeInfo (emp_name,emp_ID,designation,date_of_joining,salary,dept_name) values('jal',121,'manager','2022-03-13',40000,'cs');
... Insert into EmployeeInfo (emp_name,emp_ID,designation,date_of_joining,salary,dept_name) values('nani',131,'STAFF','2022-03-13',40000,'MARKETING');
... Insert into EmployeeInfo (emp_name,emp_ID,designation,date_of_joining,salary,dept_name) values('mont',141,'STAFF','2022-05-23',80000,'MARKETING');
... APPLY BATCH;
cqlsh:employee> SELECT * from employeeinfo
... ;

emp_id | date_of_joining | dept_name | designation | emp_name | salary
-----|-----|-----|-----|-----|-----
121 | 2022-03-12 18:30:00.000000+0000 | cs | manager | jal | 40000
141 | 2022-05-22 18:30:00.000000+0000 | MARKETING | STAFF | mont | 80000
131 | 2022-03-12 18:30:00.000000+0000 | MARKETING | STAFF | nani | 40000
(3 rows)
cqlsh:employee> update employeeinfo set emp_name='neha' where emp_id=121;
cqlsh:employee> SELECT * from employeeinfo ;

emp_id | date_of_joining | dept_name | designation | emp_name | salary
-----|-----|-----|-----|-----|-----
121 | 2022-03-12 18:30:00.000000+0000 | cs | manager | neha | 40000
141 | 2022-05-22 18:30:00.000000+0000 | MARKETING | STAFF | mont | 80000
131 | 2022-03-12 18:30:00.000000+0000 | MARKETING | STAFF | nani | 40000
(3 rows)
cqlsh:employee> update employeeinfo set emp_name='neha',dept_name='staffing' where emp_id=121;
cqlsh:employee> SELECT * from employeeinfo ;

emp_id | date_of_joining | dept_name | designation | emp_name | salary
-----|-----|-----|-----|-----|-----
121 | 2022-03-12 18:30:00.000000+0000 | staffing | manager | neha | 40000
141 | 2022-05-22 18:30:00.000000+0000 | MARKETING | STAFF | mont | 80000
131 | 2022-03-12 18:30:00.000000+0000 | MARKETING | STAFF | nani | 40000
(3 rows)
cqlsh:employee> SELECT *
... SELECT * from employeeinfo ;
cqlsh:employee> SELECT * from employeeinfo ;

emp_id | date_of_joining | dept_name | designation | emp_name | salary
-----|-----|-----|-----|-----|-----
121 | 2022-03-12 18:30:00.000000+0000 | staffing | manager | neha | 40000
141 | 2022-05-22 18:30:00.000000+0000 | MARKETING | STAFF | mont | 80000
131 | 2022-03-12 18:30:00.000000+0000 | MARKETING | STAFF | nani | 40000
(3 rows)
cqlsh:employee> SELECT * from employeeinfo order by salary;
cqlsh:employee> create index on employeeinfo(salary);
cqlsh:employee> SELECT * from employeeinfo order by salary;
cqlsh:employee> SELECT * from employeeinfo order by salary;
```

```
Activities Terminal Apr 29 10:44 bmsce@bmsce-Precision-T1700 ~
141 | 2022-05-22 18:30:00.000000+0000 | MARKETING | STAFF | mont | null | 80000
131 | 2022-03-12 18:30:00.000000+0000 | MARKETING | STAFF | nani | null | 40000
(3 rows)
cqlsh:employee> select TTL(*)
... from employeeinfo
... ;
cqlsh:employee> insert into EmployeeInfo (emp_name,emp_ID,designation,date_of_joining,salary,dept_name) values('seena',141,'STAFF','2021-05-13',50000,'MARKETING') using ttl 15;
cqlsh:employee> select ttl(*) from employeeinfo;
cqlsh:employee> select ttl(emp_name) from employeeinfo;
cqlsh:employee> select ttl(emp_name) from employeeinfo;

ttl(emp_name)
-----
null
null
(2 rows)
cqlsh:employee> select ttl(emp_name,emp_ID) from employeeinfo;
cqlsh:employee> select ttl(emp_name) from employeeinfo;

ttl(emp_name)
-----
null
null
(2 rows)
cqlsh:employee> cd ..
...
...
...
Incomplete statement at end of file
bmsce@bmsce-Precision-T1700:~$ cqlsh
Connected to test cluster at 127.0.0.1:9042.
[cqlsh 5.0.1 | Cassandra 3.11.4 | CQL spec 3.4.4 | Native protocol v4]
Use HELP for help.
cqlsh> use employee;
... ;
cqlsh:employee> select * from employeeinfo;

emp_id | date_of_joining | dept_name | designation | emp_name | projects | salary
-----|-----|-----|-----|-----|-----|-----
121 | 2022-03-12 18:30:00.000000+0000 | staffing | manager | neha | ['ai ml', 'cyber security', 'data science'] | 40000
131 | 2022-03-12 18:30:00.000000+0000 | MARKETING | STAFF | nani | null | 40000
(2 rows)
cqlsh:employee> create index on employeeinfo(emp_name);
cqlsh:employee> select * from employeeinfo where emp_name='neha';
cqlsh:employee> select * from employeeinfo where emp_name='neha';
cqlsh:employee> select * from employeeinfo where emp_name='neha';

emp_id | date_of_joining | dept_name | designation | emp_name | projects | salary
-----|-----|-----|-----|-----|-----|-----
```

2.Perform the following DB operations using Cassandra.

1.Create a keyspace by name Library

2. Create a column family by name Library-Info with

attributes Stud_Id Primary Key, Counter_value of type

Counter, Stud_Name, Book-Name, Book-Id, Date_of_issue

3. Insert the values into the table in batch

4. Display the details of the table created and increase the value of

the counter 5. Write a query to show that a student with id 112 has

taken a book “BDA” 2 times. 6. Export the created column to a csv

file

7. Import a given csv dataset from local file system into Cassandra column

family cqlsh:library> CREATE KEYSPACE library WITH replication = {'class':

'SimpleStrategy','replication_factor':1}; cqlsh:library> USE library ;

cqlsh:library> CREATE TABLE Library_info(stud_id int, stud_name text, book_name text,
book_id text, date_of_issue timestamp, counter_value counter, PRIMARY
KEY(stud_id,stud_name, book_name, book_id, date_of_issue));

cqlsh:library> BEGIN COUNTER BATCH

... UPDATE library_info set counter_value +=1 where stud_id = 111 and
stud_name = 'Manoj' and book_name = 'Operations Research' and book_id = '56TXT'
and date_of_issue = '2021-09-12';

... UPDATE library_info set counter_value +=1 where stud_id = 112 and
stud_name = 'Kamal' and book_name = 'Engineering Mathematics-3' and book_id =
'5ERW4' and date_of_issue = '2021-04-10';

... UPDATE library_info set counter_value +=1 where stud_id = 113 and
stud_name = 'Mahesh' and book_name = 'Robinson Crusoe' and book_id = '34EDC' and
date_of_issue = '2021-02-01';

... UPDATE library_info set counter_value +=1 where stud_id = 114 and
stud_name = 'Raj' and book_name = 'Engineering Drawing' and book_id = '123ER'
and date_of_issue = '2021-04-03';

... APPLY BATCH;

cqlsh:library> SELECT * FROM library_info ;

cqlsh:library> UPDATE library_info set counter_value += 1 where stud_id = 112 and stud_name = 'Kamal' and book_name = 'Engineering Mathematics-3' and book_id = '5ERW4' and date_of_issue = '2021-04-09';

cqlsh:library> SELECT * FROM library_info ;

```
cqlsh> use library;
cqlsh:library> create table library_info (
    ... .. stud_id int,
    ... .. stud_name text,
    ... .. book_id int,
    ... .. book_name text,
    ... .. date_of_issue timestamp,
```

```
cqlsh:library> update library_info set counter_value = counter_value+1 where stud_id = 112 and stud_name = 'Ram' and book_id = 200 and book_name = 'DSA' and date_of_issue = '2022-04-06';
cqlsh:library> update library_info set counter_value = counter_value+1 where stud_id = 113 and stud_name = 'sohan' and book_id = 300 and book_name = 'JAVA' and date_of_issue = '2022-04-07';
cqlsh:library> update library_info set counter_value = counter_value+1 where stud_id = 111 and stud_name = 'Raj' and book_id = 100 and book_name = 'ADA' and date_of_issue = '2022-04-05';
cqlsh:library> update library_info set counter_value = counter_value+1 where stud_id = 114 and stud_name = 'rohan' and book_id = 400 and book_name = 'UNIX' and date_of_issue = '2022-04-07';
cqlsh:library> select * from library_info ;
```

stud_id	book_id	stud_name	book_name	date_of_issue	counter_value
114	400	rohan	UNIX	2022-04-06 18:30:00.000000+0000	1
111	100	Raj	ADA	2022-04-04 18:30:00.000000+0000	1
112	200	Ram	DSA	2022-04-05 18:30:00.000000+0000	1
113	300	sohan	JAVA	2022-04-06 18:30:00.000000+0000	1

(4 rows)

```

113 | 300 | sohan | JAVA | 2022-04-06 18:30:00.000000+0000 | 1
(4 rows)
cqlsh:library> copy library_info(stud_id,stud_name,book_id,book_name,date_of_issue,counter_value ) to '/home/bmscece/Documents/library/lib.csv' with header=true;
Using 16 child processes

Starting copy of library.library_info with columns [stud_id, stud_name, book_id, book_name, date_of_issue, counter_value].
Processed: 4 rows; Rate: 122 rows/s; Avg. rate: 122 rows/s
4 rows exported to 1 files in 0.069 seconds.
cqlsh:library> copy library_info(stud_id,book_id,stud_name,book_name,date_of_issue,counter_value ) to '/home/bmscece/Documents/library/lib.csv' with header=true;
Using 16 child processes

Starting copy of library.library_info with columns [stud_id, book_id, stud_name, book_name, date_of_issue, counter_value].
Processed: 4 rows; Rate: 129 rows/s; Avg. rate: 129 rows/s
4 rows exported to 1 files in 0.066 seconds.
cqlsh:library> copy library_info(stud_id,book_id,stud_name,book_name,date_of_issue,counter_value ) from '/home/bmscece/Documents/library/lib.csv' with header = true;
Using 16 child processes

Starting copy of library.library_info with columns [stud_id, book_id, stud_name, book_name, date_of_issue, counter_value].
Processed: 4 rows; Rate: 7 rows/s; Avg. rate: 11 rows/s
4 rows imported from 1 files in 0.375 seconds (0 skipped).
cqlsh:library> select * from library_info ;

stud_id | book_id | stud_name | book_name | date_of_issue | counter_value
-----|-----|-----|-----|-----|-----
114 | 400 | rohan | UNIX | 2022-04-06 18:30:00.000000+0000 | 2
111 | 100 | Raj | ADA | 2022-04-04 18:30:00.000000+0000 | 2
112 | 200 | Ran | DSA | 2022-04-05 18:30:00.000000+0000 | 2
113 | 300 | sohan | JAVA | 2022-04-06 18:30:00.000000+0000 | 2
(4 rows)
cqlsh:library> truncate library_info;
cqlsh:library> select * from library_info ;

stud_id | book_id | stud_name | book_name | date_of_issue | counter_value
-----|-----|-----|-----|-----|-----
(0 rows)
cqlsh:library> copy library_info(stud_id,book_id,stud_name,book_name,date_of_issue,counter_value ) from '/home/bmscece/Documents/library/lib.csv' with header = true;
Using 16 child processes

Starting copy of library.library_info with columns [stud_id, book_id, stud_name, book_name, date_of_issue, counter_value].
Processed: 4 rows; Rate: 8 rows/s; Avg. rate: 11 rows/s
4 rows imported from 1 files in 0.376 seconds (0 skipped).
cqlsh:library> select * from library_info ;

stud_id | book_id | stud_name | book_name | date_of_issue | counter_value
-----|-----|-----|-----|-----|-----
114 | 400 | rohan | UNIX | 2022-04-06 18:30:00.000000+0000 | 1
111 | 100 | Raj | ADA | 2022-04-04 18:30:00.000000+0000 | 1
112 | 200 | Ran | DSA | 2022-04-05 18:30:00.000000+0000 | 1
113 | 300 | sohan | JAVA | 2022-04-06 18:30:00.000000+0000 | 1
(4 rows)
cqlsh:library> 

```

3. MongoDB- CRUD Demonstration

bmsce@bmsce-Precision-T1700:~\$ mongo

MongoDB shell version v3.6.8

connecting to: mongodb://127.0.0.1:27017

Implicit session: session { "id" :

UUID("d66acdb3-8482-417d-8b75-d65dae4b53ee") } MongoDB

server version: 3.6.8

> use Student

switched to db Student

> db.createCollection("student");

{ "ok" : 1 }

>

db.Student.insert({_id:1,StudName:"Megha",Grade:"vii",Hobbies:"InternetSurfing"}); WriteResult({ "nInserted" : 1 })

>

db.Student.update({_id:3,StudName:"Ayan",Grade:"vii"},{\$set:{Hobbies:"skating"}},{upsert:true}); WriteResult({ "nMatched" : 0, "nUpserted" : 1, "nModified" : 0, "_id" : 3 })

> db.Student.find({StudName:"Ayan"})

{ "_id" : 3, "Grade" : "vii", "StudName" : "Ayan", "Hobbies" : "skating" }

> db.Student.find({}, {StudName:1,Grade:1,_id:0});

{ "StudName" : "Megha", "Grade" : "vii" }

{ "Grade" : "vii", "StudName" : "Ayan" }

> db.Student.find({Grade:{ \$eq:'vii' }}).pretty();

{

"_id" : 1,

"StudName" : "Megha",

```

    "Grade" : "vii",
    "Hobbies" : "InternetSurfing"
  }
  { "_id" : 3, "Grade" : "vii", "StudName" : "Ayan",
    "Hobbies" : "skating" } >
  db.Student.find({Grade:{ $eq:'vii' }});
  { "_id" : 1, "StudName" : "Megha", "Grade" : "vii", "Hobbies" :
    "InternetSurfing" } { "_id" : 3, "Grade" : "vii", "StudName" :
    "Ayan", "Hobbies" : "skating" } >
  db.Student.find({Grade:{ $eq:'vii' }}).pretty();
  {
    "_id" : 1,
    "StudName" : "Megha",
    "Grade" : "vii",
    "Hobbies" : "InternetSurfing"
  }
  { "_id" : 3, "Grade" : "vii", "StudName" : "Ayan",
    "Hobbies" : "skating" } >
  db.Student.find({Hobbies:{ $in:['Chess','Skating'] }}).pr
etty(); >
  db.Student.find({Hobbies:{ $in:['Skating'] }}).pretty();
  > db.Student.find({Hobbies:{ $in:['skating'] }}).pretty();
  { "_id" : 3, "Grade" : "vii", "StudName" : "Ayan", "Hobbies" : "skating" }
  > db.Student.find({StudName:/^M/}).pretty();
  {
    "_id" : 1,
    "StudName" : "Megha",

```

```

        "Grade" : "vii",
        "Hobbies" : "InternetSurfing"
    }
> db.Student.find({StudName:/e/}).pretty();
{
    "_id" : 1,
    "StudName" : "Megha",
    "Grade" : "vii",
    "Hobbies" : "InternetSurfing"
}
> db.Student.c
ount(); 2
> db.Student.find().sort({StudName:-1}).pretty();
{
    "_id" : 1,
    "StudName" : "Megha",
    "Grade" : "vii",
    "Hobbies" : "InternetSurfing"
}
{ "_id" : 3, "Grade" : "vii", "StudName" : "Ayan",
"Hobbies" : "skating" } >
db.Student.save({StudName:"Vamsi",Grade
:"vi"}) WriteResult({ "nInserted" : 1 })
>
db.Students.update({_id:4},{ $set:{Location:"N
etwork"}}) WriteResult({ "nMatched" : 0,

```

```

    "nUpserted" : 0, "nModified" : 0 })
>
db.Students.update({_id:4},{Sunset:{Location:
"Network" }}) WriteResult({ "nMatched" : 0,
    "nUpserted" : 0, "nModified" : 0 })
>
db.Student.find({_id:1},{StudName:1,
Grade:1,_id:0}); { "StudName" :
    "Megha", "Grade" : "vii" }
> db.Student.find({Grade:{$ne:'VII'}}).pretty();
{
    "_id" : 1,
    "StudName" : "Megha",
    "Grade" : "vii",
    "Hobbies" : "InternetSurfing"
}
{ "_id" : 3, "Grade" : "vii", "StudName" : "Ayan",
    "Hobbies" : "skating" } {
    "_id" : ObjectId("6253f413e88b8c9e787b194e"),
    "StudName" : "Vamsi",
    "Grade" : "vi"
}
> db.Student.find({StudName:/s$/}).pretty();
>
db.Students.update({_id:3},{Sunset:{Location:null
    }}) WriteResult({ "nMatched" : 0, "nUpserted" :

```



```

0, "nModified" : 0 }) > db.Students.c
ount() 0
> db.Students.count({Grade:
"VII"}) 0
> db.Student.find({Grade:"VII").limit(3).pretty();
> db.Student.update({_id:3},{ $set:{Location:null}})
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
> db.Student.count({Grade:"VII"})
0
> db.Students.count({Grade:
"vii"}) 0
> db.Student.c
ount() 3
> db.Student.count({Grade:
"vii"}) 2
>
db.Student.find({Grade:"vii").
limit(3).pretty(); {
    "_id" : 1,
    "StudName" : "Megha",
    "Grade" : "vii",
    "Hobbies" : "InternetSurfing"
}
{
    "_id" : 3,
    "Grade" : "vii",
    "StudName" : "Ayan",

```

```

        "Hobbies" : "skating",
        "Location" : null
    }
}
>
db.Student.find().sort({StudName:1}).pretty(); {
  "_id" : 3,
  "Grade" :
  "vii",
  "StudName" : "Ayan",
  "Hobbies" : "skating",
  "Location" : null
}
{
  "_id" : 1,
  "StudName" : "Megha",
  "Grade" : "vii",
  "Hobbies" : "InternetSurfing"
}
{
  "_id" : ObjectId("6253f413e88b8c9e787b194e"),
  "StudName" : "Vamsi",
  "Grade" : "vi"
}
> db.Student.find().skip(2).pretty()
{
  "_id" : ObjectId("6253f413e88b8c9e787b194e"),

```

```

    "StudName" : "Vamsi",
    "Grade" : "vi"
}
> db.food.insert( { _id:1, fruits:['grapes','mango','apple']; } )
2022-04-11T15:05:51.894+0530 E QUERY [thread1] SyntaxError: missing ] after element
list @(shell):1:57 > db.food.insert({_id:1,fruits:['grapes','mango','ap
ple']}) WriteResult({ "nInserted" : 1 })
> db.food.insert({_id:2,fruits:['grapes','mango','che
rry']}) WriteResult({ "nInserted" : 1 })
> db.food.insert({_id:3,fruits:['banana','ma
ngo']}) WriteResult({ "nInserted" : 1 })
> db.food.find({fruits:['grapes','mango','apple']}).pretty();
{ "_id" : 1, "fruits" : [ "grapes", "mango", "apple" ] }
> db.food.find({'fruits.1':'grapes'})
> db.food.find({"fruits":{$size:2}})
{ "_id" : 3, "fruits" : [ "banana", "mango" ] }
> db.food.find({_id:1}, {"fruits":{$slice:2}})
{ "_id" : 1, "fruits" : [ "grapes", "mango" ] }
> db.food.find({fruits:{$all:["mango","grapes"]}})
{ "_id" : 1, "fruits" : [ "grapes", "mango", "apple" ] }
{ "_id" : 2, "fruits" : [ "grapes", "mango", "cherry" ] }
> db.food.update({_id:3},{ $set:{"fruits.1":"apple"}})
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
> db.food.update({_id:2},{ $push: { price: { grapes: 80, mango: 200, cherry: 1
00 } } }) WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })

>db.Customers.insert({_custID:1,AcctBal:'100000',AcctType:"saving"});

```

```
WriteResult({ "nInserted" : 1 })
```

```
> db.Customers.aggregate({$group:{_id:"$custID",TotAccBal:{$sum:"$AccBal"}}});
```

```
{ "_id" : null, "TotAccBal" : 0 }
```

```
db.Customers.aggregate({$match:{AcctType:"saving"}},{ $group:{_id:"$custID",TotAccBal:{$sum:"$AccBal"}}}); { "_id" : null, "TotAccBal" : 0 }
```

```
db.Customers.aggregate({$match:{AcctType:"saving"}},{ $group:{_id:"$custID",TotAccBal:{$sum:"$AccBal"}},{ $ match:{TotAccBal:{$gt:1200}}}
```

4. Screenshot of Hadoop installed

```
Microsoft Windows [Version 10.0.22000.739]
(c) Microsoft Corporation. All rights reserved.

C:\WINDOWS\system32>start-all.cmd
This script is Deprecated. Instead use start-dfs.cmd and start-yarn.cmd
starting yarn daemons
f
C:\WINDOWS\system32>jps
7072 DataNode
13492 Jps
15844 ResourceManager
16196 NameNode
1388 NodeManager

C:\WINDOWS\system32>hdfs dfs -ls -R /
drwxr-xr-x - khush supergroup 0 2022-06-27 14:09 /input
drwxr-xr-x - khush supergroup 0 2022-06-21 09:03 /input/inputtest
-rw-r--r-- 1 khush supergroup 21 2022-06-21 09:03 /input/inputtest/output.txt
-rw-r--r-- 1 khush supergroup 21 2022-06-21 08:19 /input/sample.txt
-rw-r--r-- 1 khush supergroup 21 2022-06-27 14:09 /input/sample2.txt
drwxr-xr-x - khush supergroup 0 2022-06-21 13:30 /test
-rw-r--r-- 1 khush supergroup 19 2022-06-21 13:30 /test/sample.txt

C:\WINDOWS\system32>hadoop version
Hadoop 3.3.3
Source code repository https://github.com/apache/hadoop.git -r d37506cbda38c338d9fe481addda5a05fb516f71
Compiled by stevel on 2022-05-09T16:36Z
Compiled with protoc 3.7.1
From source with checksum eb96dd4a797b6989ae0c9db66efc6
This command was run using /C:/hadoop-3.3.3/share/hadoop/common/hadoop-common-3.3.3.jar

C:\WINDOWS\system32>
```

5. Execution of HDFS Commands for interaction with Hadoop Environment.

```
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -cat /mydir/file1.txt
21/04/19 23:38:07 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
I am using Hadoop
line1
line2
```

```
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -copyFromLocal ~/file1.txt /mydir
21/04/19 23:19:36 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -ls /mydir
21/04/19 23:20:13 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
-rw-r--r--  1 hduser supergroup          30 2021-04-19 23:19 /mydir/file1.txt
hduser@lab-VirtualBox:/usr/local/sbin$
```

```
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -copyToLocal /mydir ~/hadoopcopy
21/04/19 23:29:39 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
hduser@lab-VirtualBox:/usr/local/sbin$
```

```
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -ls /
21/04/19 23:48:41 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
drwxr-xr-x  - hduser supergroup          0 2021-04-19 23:45 /mydir
drwxr-xr-x  - hduser supergroup          0 2021-04-19 23:41 /newdir
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -cp /mydir/sample.txt /newdir
21/04/19 23:48:56 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -ls /newdir
21/04/19 23:49:22 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
drwxr-xr-x  - hduser supergroup          0 2021-04-19 23:21 /newdir/mydir
-rw-r--r--  1 hduser supergroup          13 2021-04-19 23:48 /newdir/sample.txt
hduser@lab-VirtualBox:/usr/local/sbin$
```



```
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -mkdir /mydir
21/04/19 22:58:30 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
```

```
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -ls /
21/04/19 23:41:08 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 3 items
```

```
drwxr-xr-x - hduser supergroup      0 2021-04-19 23:19 /mydir
drwxr-xr-x - hduser supergroup      0 2021-04-19 23:21 /mydr
drwxr-xr-x - hduser supergroup      0 2021-04-19 23:39 /newdir
```

```
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -mv /mydr /newdir
21/04/19 23:41:38 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
```

```
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -ls /
21/04/19 23:41:44 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
```

```
drwxr-xr-x - hduser supergroup      0 2021-04-19 23:19 /mydir
drwxr-xr-x - hduser supergroup      0 2021-04-19 23:41 /newdir
```

```
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -ls /newdir
21/04/19 23:42:05 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
```

```
drwxr-xr-x - hduser supergroup      0 2021-04-19 23:21 /newdir/mydr
hduser@lab-VirtualBox:/usr/local/sbin$
```

```
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -put ~/file1.txt /mydr
21/04/19 23:21:41 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
```

```
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -ls /mydr
hadoop: command not found
```

```
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -ls /mydr
21/04/19 23:22:20 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
```

```
-rw-r--r-- 1 hduser supergroup      30 2021-04-19 23:21 /mydr/file1.txt
```

6. Create a Map Reduce program to

a) find average temperature for each year from the NCDC data set. b) find the mean max temperature for every month

AverageDriver

```
package temp;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import
org.apache.hadoop.mapreduce.lib.output.FileOutput
utFormat; public class AverageDriver {

    public static void main(String[] args) throws Exception {
        if (args.length != 2) {
            System.err.println("Please Enter the input and output parameters");
            System.exit(-1);
        }

        Job job = new Job();
        job.setJarByClass(AverageDriver.class);
        job.setJobName("Max temperature");
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
```

```

job.setMapperClass(AverageMapper.class);

job.setReducerClass(AverageReducer.class);
job.setOutputKeyClass(Text.class);

job.setOutputValueClass(IntWritable.class);

System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

AverageMapper

```

package temp;

import java.io.IOException;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

public class AverageMapper extends Mapper<LongWritable, Text,
Text, IntWritable> { public static final int MISSING = 9999;

public void map(LongWritable key, Text value,
Mapper<LongWritable, Text, Text, IntWritable>.Context
context) throws IOException, InterruptedException { int
temperature;

String line = value.toString();

String year = line.substring(15, 19);

if (line.charAt(87) == '+') {

temperature = Integer.parseInt(line.substring(88, 92));

```



```

    } else {
        temperature = Integer.parseInt(line.substring(87, 92));
    }

    String quality = line.substring(92, 93);
    if (temperature != 9999 && quality.matches("[01459]"))
        context.write(new Text(year), new IntWritable(temperature));
    }
}

```

AverageReducer

```

package temp;

import java.io.IOException;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

public class AverageReducer extends Reducer<Text, IntWritable, Text, IntWritable> {

    public void reduce(Text key, Iterable<IntWritable> values, Reducer<Text, IntWritable, Text,
IntWritable>.Context context) throws IOException, InterruptedException {

        int max_temp = 0;

        int count = 0;

        for (IntWritable value : values) {

            max_temp += value.get();

            count++;

        }

        context.write(key, new IntWritable(max_temp / count));

    }
}

```

```
}
```

```
c:\hadoop_new\sbin>hdfs dfs -cat /tempAverageOutput/part-r-00000
1901      46
1949      94
1950       3
```

b)MeanMaxDriver.class

```
package meanmax;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import
org.apache.hadoop.mapreduce.lib.input.FileInputF
ormat; import
org.apache.hadoop.mapreduce.lib.output.FileOutp
utFormat; public class MeanMaxDriver {
    public static void main(String[] args)
        throws Exception { if (args.length !=
        2) {
            System.err.println("Please Enter the input and
            output parameters"); System.exit(-1);
        }
        Job job = new Job();
        job.setJarByClass(MeanMaxDriver.class);
        job.setJobName("Max temperature");
        FileInputFormat.addInputPath(job, new
```

```

Path(args[0]));
FileOutputFormat.setOutputPath(job,
new Path(args[1]));
job.setMapperClass(MeanMaxMapper.class);
job.setReducerClass(MeanMaxReducer.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);
System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}
MeanMaxMapper.class

```

```

package meanmax;

import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

public class MeanMaxMapper extends Mapper<LongWritable, Text,
Text, IntWritable> { public static final int MISSING = 9999;

public void map(LongWritable key, Text value,
Mapper<LongWritable, Text, Text, IntWritable>.Context
context) throws IOException, InterruptedException { int
temperature;

String line = value.toString();

```

```

String month = line.substring(19, 21);
if (line.charAt(87) == '+') {
temperature = Integer.parseInt(line.substring(88, 92));
} else {
temperature = Integer.parseInt(line.substring(87, 92));
}
String quality = line.substring(92, 93);
if (temperature != 9999 && quality.matches("[01459]"))
context.write(new Text(month), new IntWritable(temperature));
}
}

```

MeanMaxReducer.class

```

package meanmax;

import java.io.IOException;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

public class MeanMaxReducer extends Reducer<Text, IntWritable, Text, IntWritable> {

    public void reduce(Text key, Iterable<IntWritable> values, Reducer<Text,
IntWritable, Text, IntWritable>.Context context) throws IOException,
InterruptedException {

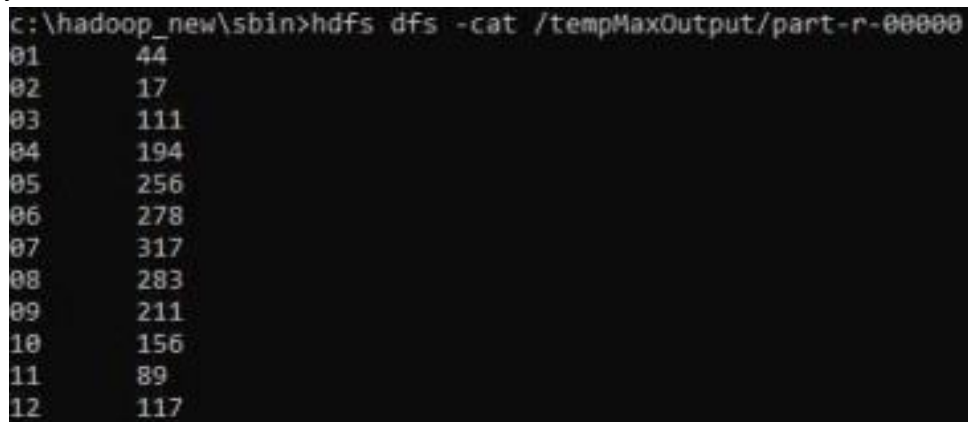
        int max_temp = 0;
        int total_temp = 0;
        int count = 0;
        int days = 0;

        for (IntWritable value : values) {
            int temp = value.get();

```

```
if (temp > max_temp)
    max_temp = temp;
    count++;
    if (count == 3) {
        total_temp += max_temp;
        max_temp = 0;
        count = 0;
        days++;
    }
}

context.write(key, new IntWritable(total_temp / days));
}
}
```



A terminal window showing the execution of an HDFS command. The command is `c:\hadoop_new\sbin>hdfs dfs -cat /tempMaxOutput/part-r-00000`. The output consists of 12 lines, each with a two-digit number followed by a space and a three-digit number. The numbers are: 01 44, 02 17, 03 111, 04 194, 05 256, 06 278, 07 317, 08 283, 09 211, 10 156, 11 89, and 12 117.

Index	Value
01	44
02	17
03	111
04	194
05	256
06	278
07	317
08	283
09	211
10	156
11	89
12	117

7. For a given Text file, Create a Map Reduce program to sort the content in an alphabetic order listing only top 10 maximum occurrences of words.

```
//Driver Code

package wordCount;

import java.io.IOException;

import org.apache.hadoop.conf.Configured;

import org.apache.hadoop.fs.Path;

import org.apache.hadoop.io.IntWritable;

import org.apache.hadoop.io.Text;

import org.apache.hadoop.mapred.FileInputFormat;

import org.apache.hadoop.mapred.FileOutputFormat;


import org.apache.hadoop.mapred.JobClient;

import org.apache.hadoop.mapred.JobConf;

import org.apache.hadoop.util.Tool;

import org.apache.hadoop.util.ToolRunner;


public class WCDriver extends Configured implements Tool {

    public int run(String args[]) throws IOException
    {

        if (args.length < 2)

        {

            System.out.println("Please give valid inputs");

            return -1;

        }

    }

}
```

```

        JobConf conf = new JobConf(WCDriver.class);
        FileInputFormat.setInputPaths(conf, new
        Path(args[0]));
        FileOutputFormat.setOutputPath(conf,
        new Path(args[1]));
        conf.setMapperClass(WCMapper.class);
        conf.setReducerClass(WCReducer.class);
        conf.setMapOutputKeyClass(Text.class);
        conf.setMapOutputValueClass(IntWritable.class);
        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(IntWritable.class);
        JobClient.runJob(conf);
        return 0;
    }
    // Main Method
    public static void main(String args[]) throws Exception
    {
        int exitCode = ToolRunner.run(new WCDriver(), args);
        System.out.println(exitCode);
    }
}

```

//Mapper Code

```

package wordCount;
import java.io.IOException;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;

```

```

import org.apache.hadoop.io.Text;

import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.OutputCollector;


import org.apache.hadoop.mapred.Reporter;

public class WCMapper extends MapReduceBase implements Mapper<LongWritable,Text,
    Text, IntWritable> { // Map function

    public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable>
output, Reporter rep) throws IOException
    {
        String line = value.toString();
        // Splitting the line on spaces
        for (String word : line.split(" "))
        {
            if (word.length() > 0)
            {
                output.collect(new Text(word), new IntWritable(1));
            }
        }
    }
}

```

```

//Reducer Code
package wordCount;

```

```

import java.io.IOException;

```



```
import java.util.Iterator;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;

public class WCReducer extends MapReduceBase implements Reducer<Text,IntWritable,
    Text, IntWritable> { // Reduce function

    public void reduce(Text key, Iterator<IntWritable> value,
OutputCollector<Text, IntWritable> output,Reporter rep) throws IOException
    {

        int count = 0;

        // Counting the frequency of each words
        while (value.hasNext())
        {

            IntWritable i = value.next();

            count += i.get();

        }

        output.collect(key, new IntWritable(count));

    }

}
```

```
Activities Terminal Apr 26 15:26
hduser@lab-VirtualBox: /home/lab/hadoop-2.6.0/share/ha...

Bytes Read=44
File Output Format Counters
Bytes Written=35
hduser@lab-VirtualBox:/home/lab/hadoop-2.6.0/share/hadoop/mapreduce$ hadoop dfs
-ls /firstExampleOut
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.

21/04/26 15:20:15 WARN util.NativeCodeLoader: Unable to load native-hadoop libr
ary for your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r-- 1 hduser supergroup 0 2021-04-26 15:19 /firstExampleOut/_
SUCCESS
-rw-r--r-- 1 hduser supergroup 35 2021-04-26 15:19 /firstExampleOut/p
art-r-000000
hduser@lab-VirtualBox:/home/lab/hadoop-2.6.0/share/hadoop/mapreduce$ hadoop dfs
-cat /firstExampleOut/part-r-000000
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.

21/04/26 15:22:01 WARN util.NativeCodeLoader: Unable to load native-hadoop libr
ary for your platform... using builtin-java classes where applicable
bear 2
car 3
deer 1
deer 1
river 2
hduser@lab-VirtualBox:/home/lab/hadoop-2.6.0/share/hadoop/mapreduce$ hadoop dfs
-cat /firstExampleOut/part-r-000000
```

8. Create a Map Reduce program to demonstrating join operation

```
// JoinDriver.java

import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.*;
import org.apache.hadoop.mapred.lib.MultipleInputs;
import org.apache.hadoop.util.*;

public class JoinDriver extends Configured implements Tool {

    public static class KeyPartitioner implements Partitioner<TextPair, Text> {

        @Override

        public void configure(JobConf job) {}

        @Override

        public int getPartition(TextPair key, Text value, int numPartitions) {

            return (key.getFirst().hashCode() & Integer.MAX_VALUE) %

            numPartitions;

        }

    }

    @Override

    public int run(String[] args) throws Exception {

        if (args.length != 3) {

            System.out.println("Usage: <Department Emp Strength input>

            <Department Name input>

            <output>"); return -1;

        }

    }

}
```

```

}
JobConf conf = new JobConf(getConf(), getClass());

conf.setJobName("Join 'Department Emp Strength input' with 'Department Nameinput'");

Path AInputPath = new Path(args[0]);

Path BInputPath = new Path(args[1]);

Path outputPath = new Path(args[2]);

MultipleInputs.addInputPath(conf, AInputPath,
    TextInputFormat.class, Posts.class);

MultipleInputs.addInputPath(conf, BInputPath,
    TextInputFormat.class, User.class);

FileOutputFormat.setOutputPath(conf, outputPath);

conf.setPartitionerClass(KeyPartitioner.class);

conf.setOutputValueGroupingComparator(TextPair.FirstComparator.class);

conf.setMapOutputKeyClass(TextPair.class);

conf.setReducerClass(JoinReducer.class); conf.setOutputKeyClass(Text.class);

JobClient.runJob(conf);

return 0;
}

public static void main(String[] args) throws Exception {

    int exitCode = ToolRunner.run(new JoinDriver(), args);

    System.exit(exitCode);

}

}

```

// JoinReducer.java

```
import java.io.IOException;
```

```

import java.util.Iterator;
import org.apache.hadoop.io.Text;

import org.apache.hadoop.mapred.*;

public class JoinReducer extends MapReduceBase implements Reducer<TextPair, Text,
Text,
Text> {

@Override

public void reduce (TextPair key, Iterator<Text> values, OutputCollector<Text, Text>output, Reporter reporter)
throws IOException

{

Text nodeId = new Text(values.next());
while (values.hasNext()) {
Text node = values.next();
Text outValue = new Text(nodeId.toString() + "\t\t" + node.toString());
output.collect(key.getFirst(), outValue);
}
}
}
}

```

// User.java

```

import java.io.IOException;

import java.util.Iterator;

import org.apache.hadoop.conf.Configuration;

import org.apache.hadoop.fs.FSDataInputStream;

import org.apache.hadoop.fs.FSDataOutputStream;

import org.apache.hadoop.fs.FileSystem;

```

```

import org.apache.hadoop.fs.Path;

import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text; import

org.apache.hadoop.mapred.*; import

org.apache.hadoop.io.IntWritable;

public class User extends MapReduceBase implements Mapper<LongWritable, Text,
TextPair,
Text> {

@Override

public void map(LongWritable key, Text value, OutputCollector<TextPair, Text> output,
Reporter reporter)

throws IOException

{

String valueString = value.toString();

String[] SingleNodeData = valueString.split("\t");

output.collect(new TextPair(SingleNodeData[0], "1"), new

Text(SingleNodeData[1]));

}

}

```

//Posts.java

```

import java.io.IOException;

import org.apache.hadoop.io.*;

import org.apache.hadoop.mapred.*;

public class Posts extends MapReduceBase implements Mapper<LongWritable, Text,
TextPair,
Text> {

@Override

```

```

public void map(LongWritable key, Text value, OutputCollector<TextPair, Text> output,
Reporter reporter)

throws IOException
{
String valueString = value.toString();

String[] SingleNodeData = valueString.split("\t");

output.collect(new TextPair(SingleNodeData[3], "0"), new
Text(SingleNodeData[9]));

}

}

```

// TextPair.java

```

import java.io.*;

import org.apache.hadoop.io.*;

public class TextPair implements WritableComparable<TextPair> {

private Text first;

private Text second;

public TextPair() {

set(new Text(), new Text());

}

public TextPair(String first, String second) {

set(new Text(first), new Text(second));

}

public TextPair(Text first, Text second) {

set(first, second);

}

public void set(Text first, Text second) {

this.first = first;

```

```

this.second = second;

}

public Text getFirst() {
return first;
}

public Text getSecond() {
return second;
}

@Override

public void write(DataOutput out) throws IOException {
first.write(out);

second.write(out);
}

@Override

public void readFields(DataInput in) throws IOException {
first.readFields(in);

second.readFields(in);
}

@Override

public int hashCode() {
return first.hashCode() * 163 + second.hashCode();
}

@Override

public boolean equals(Object o) {
if (o instanceof TextPair) {
TextPair tp = (TextPair) o;
return first.equals(tp.first) && second.equals(tp.second);
}
}

```



```

}

return false;

}

@Override

public String toString() {
return first + "\t" + second;

}

@Override

public int compareTo(TextPair tp) {
int cmp = first.compareTo(tp.first);
if (cmp != 0) {
return cmp;
}

return second.compareTo(tp.second);
}

// ^^ TextPair

// vv TextPairComparator

public static class Comparator extends WritableComparator {
private static final Text.Comparator TEXT_COMPARATOR = new Text.Comparator();

public Comparator() {
super(TextPair.class);
}

@Override

public int compare(byte[] b1, int s1, int l1,
byte[] b2, int s2, int l2) {
try {

int firstL1 = WritableUtils.decodeVIntSize(b1[s1]) + readVInt(b1, s1);

```

```

int firstL2 = WritableUtils.decodeVIntSize(b2[s2]) + readVInt(b2, s2);
int cmp = TEXT_COMPARATOR.compare(b1, s1, firstL1, b2, s2, firstL2);
if (cmp != 0) {
    return cmp;
}

return TEXT_COMPARATOR.compare(b1, s1 + firstL1, l1 - firstL1,
b2, s2 + firstL2, l2 - firstL2);
} catch (IOException e) {
    throw new IllegalArgumentException(e);
}
}
}

static {
    WritableComparator.define(TextPair.class, new Comparator());
}

public static class FirstComparator extends WritableComparator { private static
final Text.Comparator TEXT_COMPARATOR = new Text.Comparator(); public
FirstComparator() {
    super(TextPair.class);
}

@Override
public int compare(byte[] b1, int s1, int l1,
byte[] b2, int s2, int l2) {
    try {
        int firstL1 = WritableUtils.decodeVIntSize(b1[s1]) + readVInt(b1, s1);
        int firstL2 = WritableUtils.decodeVIntSize(b2[s2]) + readVInt(b2, s2);
        return TEXT_COMPARATOR.compare(b1, s1, firstL1, b2, s2, firstL2);
    }
}

```

```

    } catch (IOException e) {
        throw new IllegalArgumentException(e);
    }
}

@Override
public int compare(WritableComparable a, WritableComparable b) {
    if (a instanceof TextPair && b instanceof TextPair) {
        return (((TextPair) a).first.compareTo(((TextPair) b).first));
    }
    return super.compare(a, b);
}
}}

```

```

hduser@bmsce-Precision-T1700:/home/bmsce$ hdfs dfs -cat /join/output/*
A11      Finance      50
B12      HR              100
C13      Manufacturing    250
Dept_ID  Dept_Name        Total_Employee

```

9. Program to print word count on scala shell and print “Hello world” on scala IDE

```
val data=sc.textFile("sparkdata.txt")

data.collect;

val splitdata = data.flatMap(line => line.split(" "));

splitdata.collect;

val mapdata = splitdata.map(word => (word,1));

mapdata.collect;

val reducedata = mapdata.reduceByKey(_+_);

reducedata.collect;
```

A screenshot of a Scala REPL window titled "binase@bssuz-Pradison-f1780: ~/Desktop". The window shows the Spark logo and version 2.4.0. The user enters several Scala commands to read a file, split it into words, map each word to a count of 1, and then reduce by key to get the total count for each word. The output shows the raw text of the file, the split words, the mapped pairs, and the final reduced word counts.

```
spark sessions available as 'spark',
welcome to

Spark version 2.4.0

Using Scala version 2.11.12 (OpenJDK 64-Bit Server VM, Java 1.8.0_212)
Type in expressions to have them evaluated.
Type :help for more information.

scala> val data=sc.textFile("sample.txt")
data: org.apache.spark.rdd.RDD[String] = sample.txt MapPartitionsRDD[1] at textFile at <console>:14

scala> data.collect;
res0: Array[String] = Array(hi how are you, how is your job, how is your family, how is your brother, how is your sister)

scala> val splitdata = data.flatMap(line => line.split(" "));
splitdata: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[2] at FlatMap at <console>:25

scala> splitdata.collect;
res1: Array[String] = Array(hi, how, are, you, how, is, your, job, how, is, your, family, how, is, your, brother, how, is, your, sister)

scala> val mapdata = splitdata.map(word => (word,1));
mapdata: org.apache.spark.rdd.RDD[(String, Int)] = MapPartitionsRDD[3] at map at <console>:28

scala> mapdata.collect;
res2: Array[(String, Int)] = Array((hi,1), (how,1), (are,1), (you,1), (how,1), (is,1), (your,1), (job,1), (how,1), (is,1), (your,1), (family,1), (how,1), (is,1), (your,1), (brother,1), (how,1), (is,1), (your,1), (sister,1))

scala> val reducedata = mapdata.reduceByKey(_+_);
reducedata: org.apache.spark.rdd.RDD[(String, Int)] = ShuffleRDD[4] at reduceByKey at <console>:29

scala> reducedata.collect;
res3: Array[(String, Int)] = Array((are,1), (brother,1), (is,4), (sister,1), (family,1), (how,5), (job,1), (you,1), (hi,1), (your,4))

scala>
```

10. Using RDD and FlMap count how many times each word appears in a file and write out a list of words whose count is strictly greater than 4 using Spark

```
val textFile = sc.textFile("/home/bhoom/Desktop/wc.txt")
val counts = textFile.flatMap(line => line.split(" ")).map(word => (word, 1)).reduceByKey(_ + _)
import scala.collection.immutable.ListMap
val sorted=ListMap(counts.collect.sortWith(_._2 > _._2):_*)// sort in descending order based
on values
println(sorted)
for((k,v)<-sorted)
{
  if(v>4)
  {
    print(k+",")
    print(v)
    println()
  }
}
```

```

scala> val textFile = sc.textFile("sample.txt")
textFile: org.apache.spark.rdd.RDD[String] = sample.txt MapPartitionsRDD[1] at textFile at <console>:24

scala> val counts = textFile.flatMap(line => line.split(" ")).map(word => (word, 1)).reduceByKey(_ + _)
counts: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[4] at reduceByKey at <console>:25

scala> import scala.collection.immutable.ListMap
import scala.collection.immutable.ListMap

scala> val sorted=ListMap(counts.collect.sortWith(_._2 > _._2):_*)// sort in descending order based
sorted: scala.collection.immutable.ListMap[String,Int] = Map(how -> 5, is -> 4, your -> 4, are -> 1, brother -> 1, sister -> 1, family -> 1, job -> 1, you -> 1, hl -> 1)

scala> println(sorted)
Map(how -> 5, is -> 4, your -> 4, are -> 1, brother -> 1, sister -> 1, family -> 1, job -> 1, you -> 1, hl -> 1)

scala> for((k,v)<-sorted)
| {
|   |
Display all 689 possibilities? (y or n)
|   |
|   | if(v>4)
|   | {
|   |   | print(k+",")
|   |   | print(v)
|   |   | println()
|   |   | }
|   | }
| }
how,5

```